

Multiple Linear Regression

```
In [1]: #Import libraries
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
%matplotlib inline

In [2]: # Import data file

dataset=pd.read_csv('insurance.csv')

dataset.head()
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

Encoding categorical data

```
In [3]: objList = dataset.select_dtypes(include = "object").columns
print (objList)

Index(['sex', 'smoker', 'region'], dtype='object')

In [4]: #Label Encoding for object to numeric conversion
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

for feat in objList:
    dataset[feat] = le.fit_transform(dataset[feat].astype(str))

print (dataset.info())

dataset.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   int64
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   int64
5   region      1338 non-null   int64
6   charges     1338 non-null   float64
dtypes: float64(2), int64(5)
memory usage: 73.3 KB
None
```

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520

Create X and y arrays

```
In [5]: # Create X set and y set. They are arrays and vectors
#X=dataset.iloc[:,0].values
X=dataset.iloc[:, :-1].values
print(X)
y=dataset.iloc[:, -1].values

print(y)
```

```
[[19.  0.  27.9  0.  1.  3. ]
 [18.  1.  33.77 1.  0.  2. ]
 [28.  1.  33.   3.  0.  2. ]
 ...
 [18.  0.  36.85 0.  0.  2. ]
 [21.  0.  25.8  0.  0.  3. ]
 [61.  0.  29.07 0.  1.  1. ]]
[16884.924 1725.5523 4449.462 ... 1629.8335 2007.945 29141.3603]
```

Creating Training and Test Data Sets

```
In [6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test=train_test_split(X,y, test_size=0.2, random_state=44)
```

Training Multiple Linear Regression on Training set

```
In [7]: from sklearn.linear_model import LinearRegression #Calling LinearRegression Class
regressor=LinearRegression() #Creating an object of LinearRegression
regressor.fit(X_train, y_train) #Method of the linear regression class

#Note: LinearRegression Class takes care of Dummy Variable trap

Out[7]: LinearRegression()
```

Note: Backward elimination is irrelevant because the Scikit-Learn library automatically takes care of selecting the statistically significant features when training the model to make accurate predictions.

Predicting Test set results

```
In [8]: y_pred=regressor.predict(X_test) # Predicting values on test data
#print(y_pred)

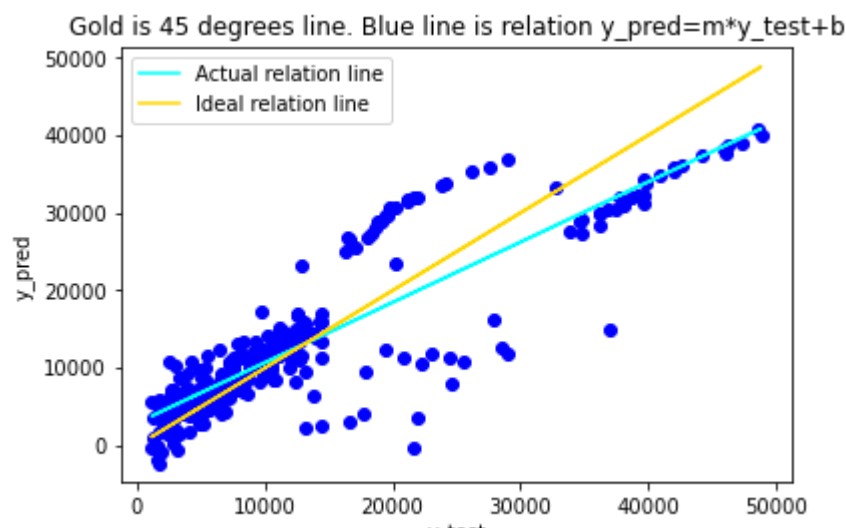
In [9]: np.set_printoptions(precision=2)
# To display with two decimal points all printed values

In [10]: #print(np.concatenate((y_pred.reshape(len(y_pred),1),
#y_test.reshape(len(y_test), 1)),axis=1))
#Concatenate both y_pred amd y_test as column vectors
```

Representation of Y_pred vs Y_test in simple scatter plot and line representation.

The gold line represents an ideal match between "measuring devices" The blue line represents the actual match between Y_pred and Y_test.

```
In [11]: plt.scatter(y_test, y_pred, color='blue')
m, b = np.polyfit(y_test, y_pred, 1) # 1 is the degree of the polynomial to fit.
plt.plot(y_test, m*y_test+b, color='cyan', label="Actual relation line")
plt.plot(y_test, y_test, color='gold', label="Ideal relation line")
plt.title("Gold is 45 degrees line. Blue line is relation y_pred=m*y_test+b")
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.legend()
plt.show()
print("Slope: m=", "{:.2f}".format(m))
print("Intercept: b=", "{:.2f}".format(b))
```



Slope: m= 0.78
Intercept: b= 2924.65

Calculating Regression model's precision (Aspect Ratio: AR)

```
In [12]: #Estimating correlation coefficient and R^2 between Y_pred and Y_test

correlation_matrix = np.corrcoef(y_test, y_pred)
correlation_xy = correlation_matrix[0,1]
r_squared = correlation_xy**2
print("R-squared between y_pred and y_test: ", "{:.2f}".format(r_squared))

R-squared between y_pred and y_test:  0.75

In [13]: # Aspect ratio calculation (precision)
AR=(1-correlation_xy)**(0.5-1)
print(" 'Aspect Ratio': ", "{:.2f}".format(AR))

'Aspect Ratio':  2.76
```

A good correlation coefficient is >=0.98, yielding an aspect ratio of 7 to 1 (minimum suggested when comparing a measuring device (y_test or the regression model) against a master measuring device (y_test)).

Obtaining model coefficients

```
In [14]: # Obtaining the model equation
print(regressor.coef_)
print(regressor.intercept_)

[ 254.21 -233.99 345.45 502.46 23966.6 -344.46]
-12095.637432727546
```

Explained variance score for Multiple Regression

The Explained Variance Score (EVS) best possible value is 1.0. Lower values are worse.

$$EVS = 1 - \frac{var(y_{test}-y_{pred})}{var(y_{test})}$$

```
In [15]: from sklearn.metrics import explained_variance_score
EVS=explained_variance_score(y_test, y_pred, multioutput='uniform_average')
print("EVS: ", "{:.2f}".format(EVS))

EVS:  0.75
```